


Image Pre-Processing & Augmentation

What is Image Augmentation?



What do we need Image Processing?

- Need for large datasets
- Prevent Overfitting
- Feature Extraction
- Class Imbalances
- Stability of Network

Types of Augmentation

Offline Augmentation

- Used on smaller datasets
- Used to create more data from original data
- Done prior to training
- Transformations done in the beginning

Online Augmentation

- Used when dataset is large
- Used to randomly modify existing data
- Done during training
- Transformations done on minibatches in dataloader

```
class ImageData(Dataset):
    def __init__(self, df, data_dir, transform=None):
        self.df = df
        self.data_dir = data_dir
        self.transform = transform

    def __len__(self):
        return len(self.df)

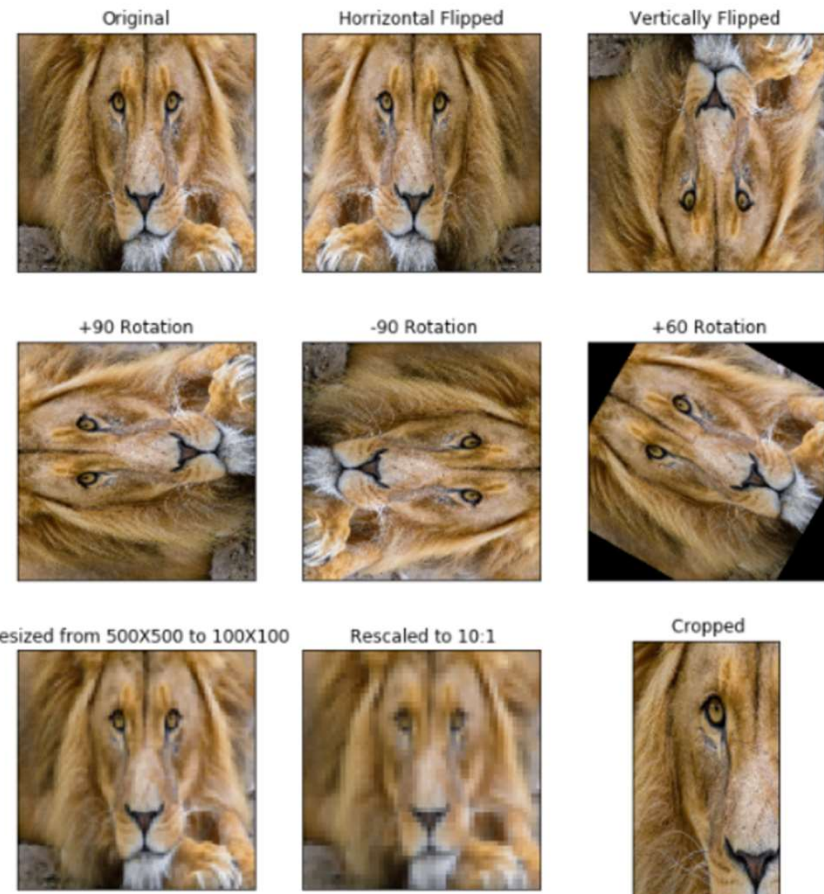
    def __getitem__(self, index):
        img_path = os.path.join(self.data_dir, self.df.iloc[index, 0])
        labels = torch.tensor(self.df.iloc[index, 1:], dtype=torch.long)
        img = Image.open(img_path).convert('RGB')
        if self.transform:
            image = self.transform(img)
        return (image, labels)
```

Online Augmentation code example

Types of Transformation

Geometric transformations


1. Flipping
2. Cropping
3. Scaling
4. Translating
5. Rotating
6. Adding Noise
7. Warping/Distorting



Types of Transformations

Colour Space Transformations

1. Colour shifting
2. Contrast
3. Sharpness
4. Blurring
5. Brightness

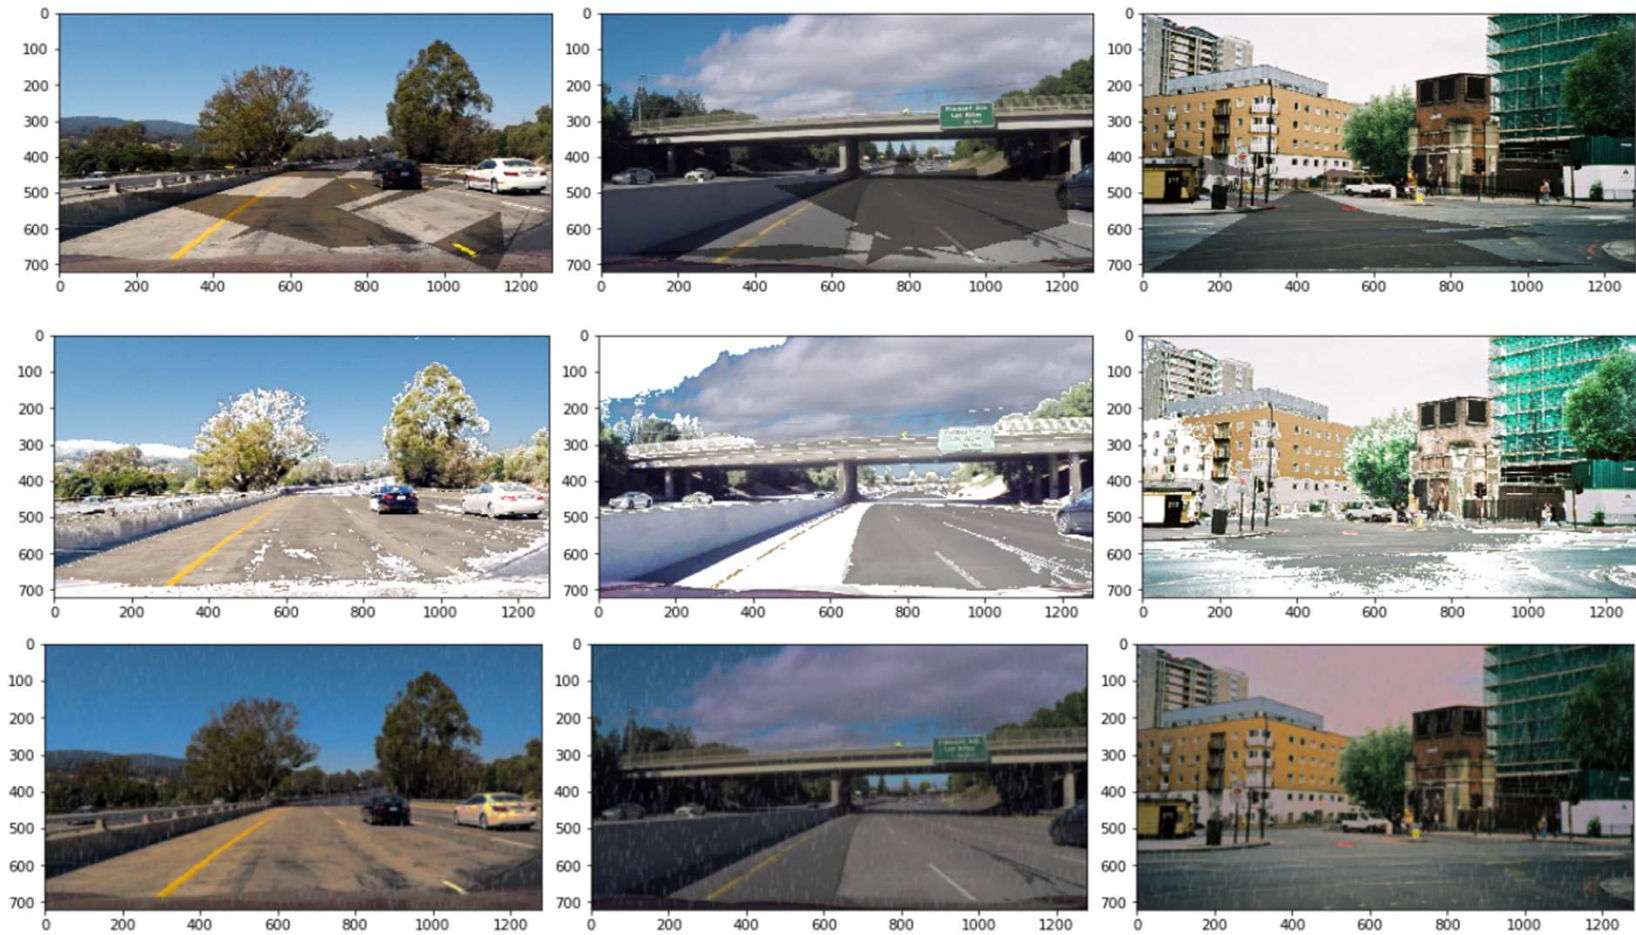







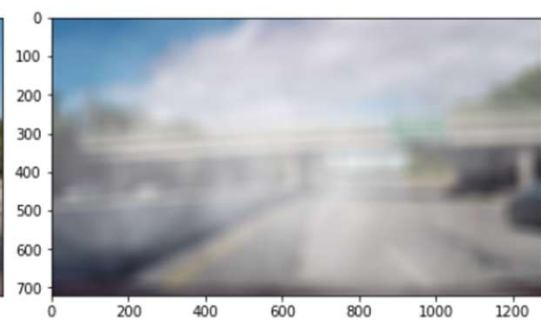
Types of Transformations

Other Techniques

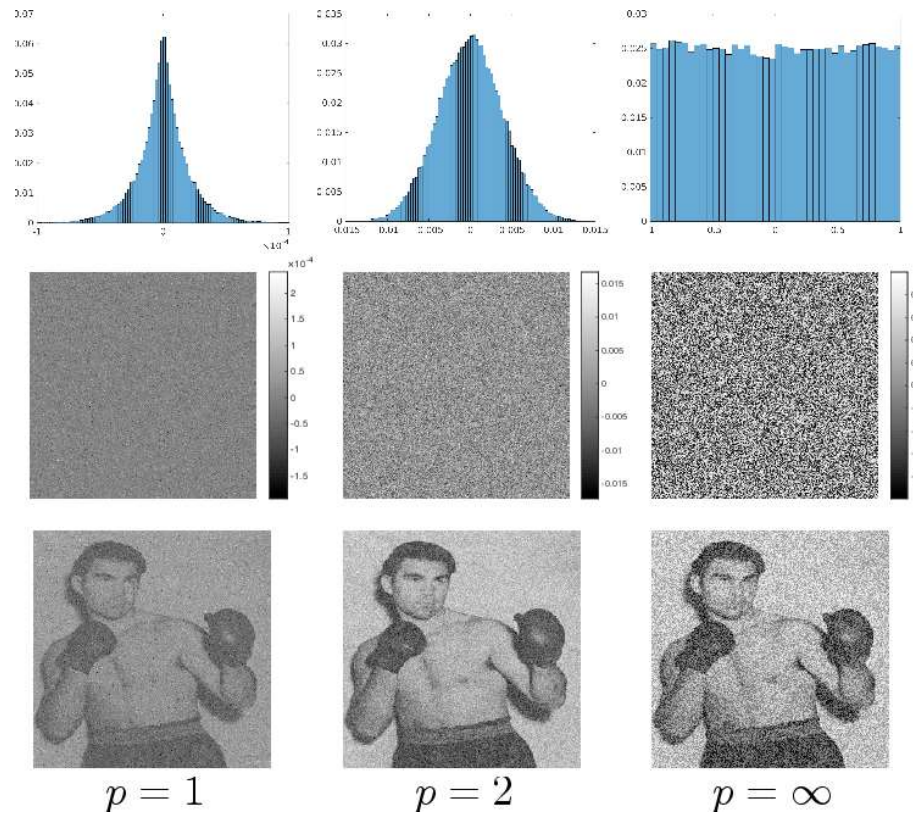
1. Pixel Dropout
2. Jitter
3. Adding
 1. Rain
 2. Snow
 3. Sunflare
 4. Shadow

Albumentation (Automold)





Gaussian Noise



Understanding effect of some fundamental transformations

Visualizing and Understanding Convolutional Networks

Matthew D. Zeiler

Dept. of Computer Science, Courant Institute, New York University

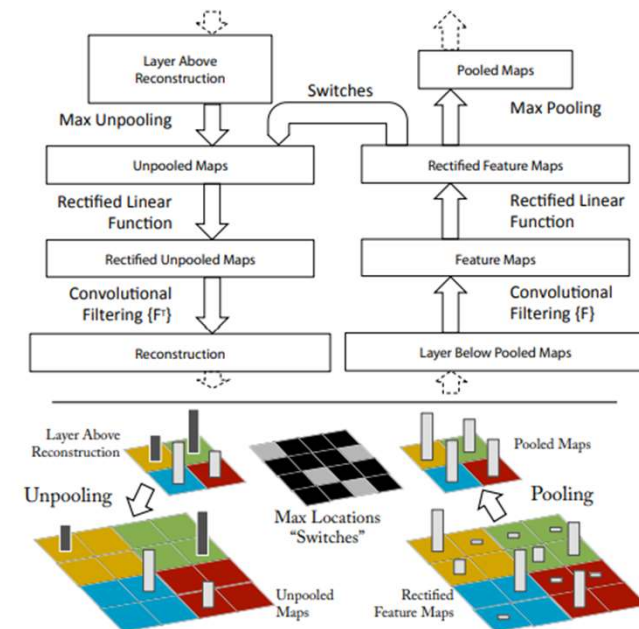
ZEILER@CS.NYU.EDU

Rob Fergus

Dept. of Computer Science, Courant Institute, New York University

FERGUS@CS.NYU.EDU

Link: <https://arxiv.org/abs/1311.2901>



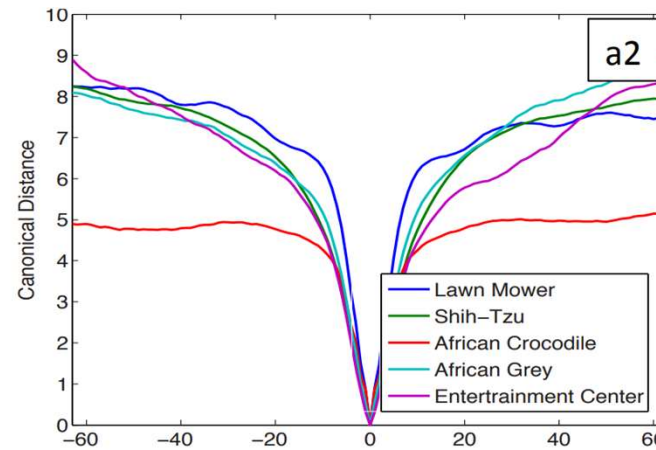



Feature Invariances




- Translation
- Scaling
- Rotation







TRANSLATION



SCALING



ROTATION



OBSERVATIONS

- 1st layer shows dramatic difference in output for any transformation
- 7th layer has lesser impact and is quasi-linear for translation and scaling
- Rotations are not corrected in the last layer unless the object is rotationally symmetric

CONCLUSION

Network output is stable to translation and scaling but is not invariant to rotation

Albumentation



Albumentations is a fast image augmentation library and easy to use wrapper around other libraries.

Albumentations package is written based on numpy, OpenCV, and imgaug. It is a very popular package written by Kaggle masters and used widely in Kaggle competitions.

Features

Albumentations package is capable of:

- Over 60 pixel-level and spatial-level transformations;
- Transforming images with masks, bounding boxes, and keypoints;
- The library is faster than other libraries on most of the transformations;
- Organizing augmentations into pipelines;
- PyTorch integration;
- Was used to get top results in many DL competitions at Kaggle, topcoder, CVPR, MICCAI;
- Written by Kaggle Masters.

Pixel-Level Transforms

Pixel-level transforms will change just an input image and will leave any additional targets such as masks, bounding boxes, and keypoints unchanged. The list of pixel-level transforms:

- Blur
- CLAHE
- ChannelShuffle
- HueSaturationValue
- RGBShift
- RandomBrightnessContrast
- ToGray
- GaussNoise


Original image




RGBShift



HueSaturationValue




ChannelShuffle



CLAHE



RandomContrast



RandomGamma



RandomBrightness



Blur




MedianBlur



ToGray




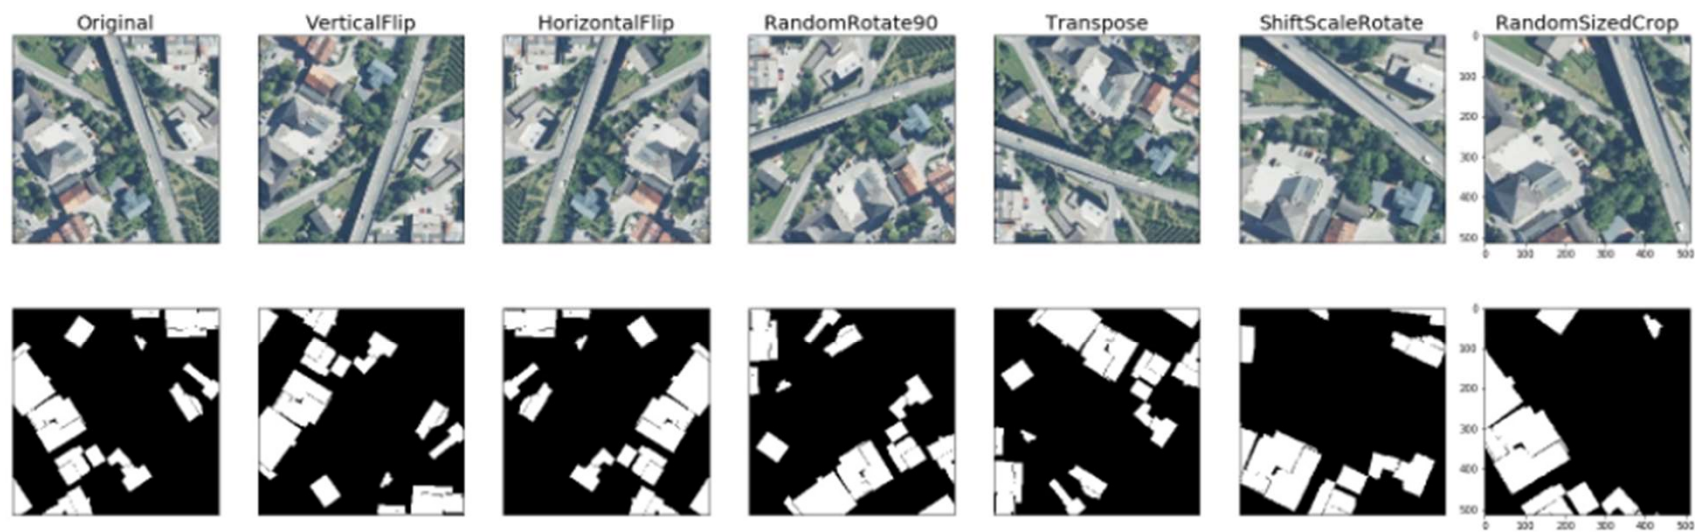
JpegCompression



Spatial-level Transforms

Spatial-level transforms will simultaneously change both an input image as well as additional targets such as masks, bounding boxes, and keypoints.

- VerticalFlip
- HorizontalFlip
- RandomRotate90
- Transpose
- RandomResizedCrop
- GridDistortion
- ElasticTransform




Original




HorizontalFlip



VerticalFlip



ShiftScaleRotate



Original image



Augmented image



Original mask



Augmented mask



Pipelines

We want to stack many transformations together into a single pipeline. Depending on the framework we are using there are different methods.

```
>>> transforms.Compose([  
>>>     transforms.CenterCrop(10),  
>>>     transforms.ToTensor(),  
>>> ])
```

```
1 # Compose a complex augmentation pipeline
2 augmentation_pipeline = A.Compose(
3     [
4         A.HorizontalFlip(p = 0.5), # apply horizontal flip to 50% of images
5         A.OneOf(
6             [
7                 # apply one of transforms to 50% of images
8                 A.RandomContrast(), # apply random contrast
9                 A.RandomGamma(), # apply random gamma
10                A.RandomBrightness(), # apply random brightness
11            ],
12            p = 0.5
13        ),
14        A.OneOf(
15            [
16                # apply one of transforms to 50% images
17                A.ElasticTransform(
18                    alpha = 120,
19                    sigma = 120 * 0.05,
20                    alpha_affine = 120 * 0.03
21                ),
22                A.GridDistortion(),
23                A.OpticalDistortion(
24                    distort_limit = 2,
25                    shift_limit = 0.5
26                ),
27            ],
28            p = 0.5
29        )
30    ],
31    p = 1
32 )
```

```

1 # Import pytorch utilities from albumentations
2 from albumentations.pytorch import ToTensor
3
4 # Define the augmentation pipeline
5 augmentation_pipeline = A.Compose(
6     [
7         A.HorizontalFlip(p = 0.5), # apply horizontal flip to 50% of images
8         A.OneOf(
9             [
10                 # apply one of transforms to 50% of images
11                 A.RandomContrast(), # apply random contrast
12                 A.RandomGamma(), # apply random gamma
13                 A.RandomBrightness(), # apply random brightness
14             ],
15             p = 0.5
16         ),
17
18         A.Normalize(
19             mean=[0.485, 0.456, 0.406],
20             std=[0.229, 0.224, 0.225]),
21
22         ToTensor() # convert the image to PyTorch tensor
23     ],
24     p = 1
25 )
26
27 # Load the augmented data
28

```

```

29 # Define the demo dataset
30 class DogDataset2(Dataset):
31     '''
32     Sample dataset for Albumentations demonstration.
33     The dataset will consist of just one sample image.
34     '''
35
36     def __init__(self, image, augmentations = None):
37         self.image = image
38         self.augmentations = augmentations # save the augmentations
39
40     def __len__(self):
41         return 1 # return 1 as we have only one image
42
43     def __getitem__(self, idx):
44         # return the augmented image
45         # no need to convert to tensor, because image is converted to tensor already by the p
46         augmented = self.augmentations(image = self.image)
47         return augmented['image']
48
49 # Initialize the dataset, pass the augmentation pipeline as an argument to init function
50 train_ds = DogDataset2(image, augmentations = augmentation_pipeline)
51
52 # Initilize the dataloader
53 trainloader = DataLoader(train_ds, batch_size=1, shuffle=True, num_workers=0)

```

Data Augmentation as a subfield

Fig. 2

